

# LAPP/SELinux

A secure web application stack using SE-PostgreSQL

KaiGai Kohei <kaigai@ak.jp.nec.com>

NEC OSS Promotion Center




# Self Introduction

**SELECT \* FROM pg\_developers WHERE name = 'KaiGai'**

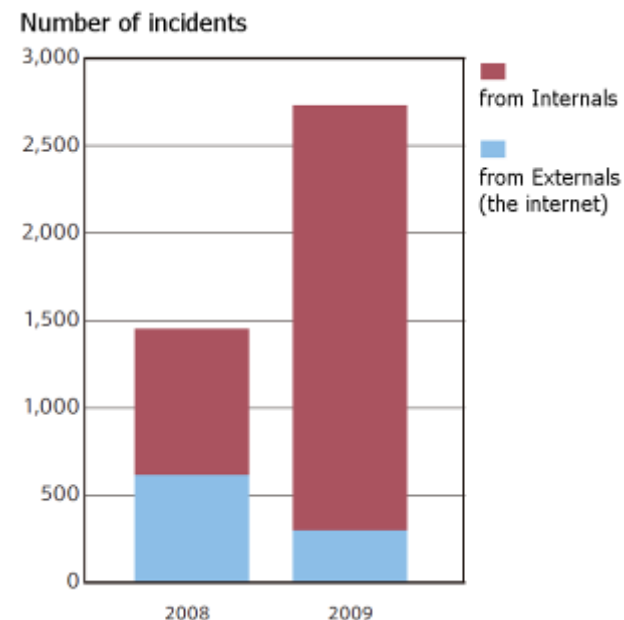

- Job NEC OSS Promotion Center, for 7 years
- Contributions
  - SMP Scalability Improvement of SELinux
  - Lead project to port SELinux into embedded platform
  - Development of SE-PostgreSQL
  - Access control support of large object, and so on...
- Interest Web system's security



# Agenda

1. Background
  2. SE-PostgreSQL
  3. Apache/SELinux plus
  4. Demonstration
  5. Future Plans
- 

# Security nightmare in web systems



(Reference: JSOC analysis report of the incursion trend, vol.12, vol.14, LAC)

- Rapid increasing of attacks to web systems
- More threats from Internals, rather than Externals
- ➡ What technology can improve the situation?

# LAPP - A typical web application stack

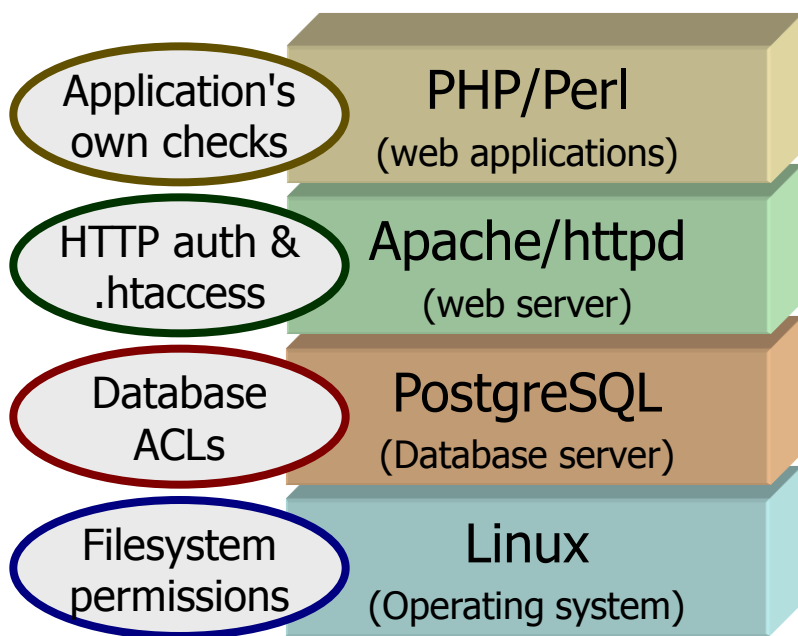
## LAPP

- Linux, Apache, PostgreSQL, PHP/Perl

## Concerns in security

- Each layer has its own security mechanism
- Web-users are not mapped to users in OS/DB

An information asset in DB being invisible might be visible in Filesystem



OS/DB layer could not distinguish actual users, so all the security burdens are pushed to web-app's

Software stack like a symphony,  
But ....





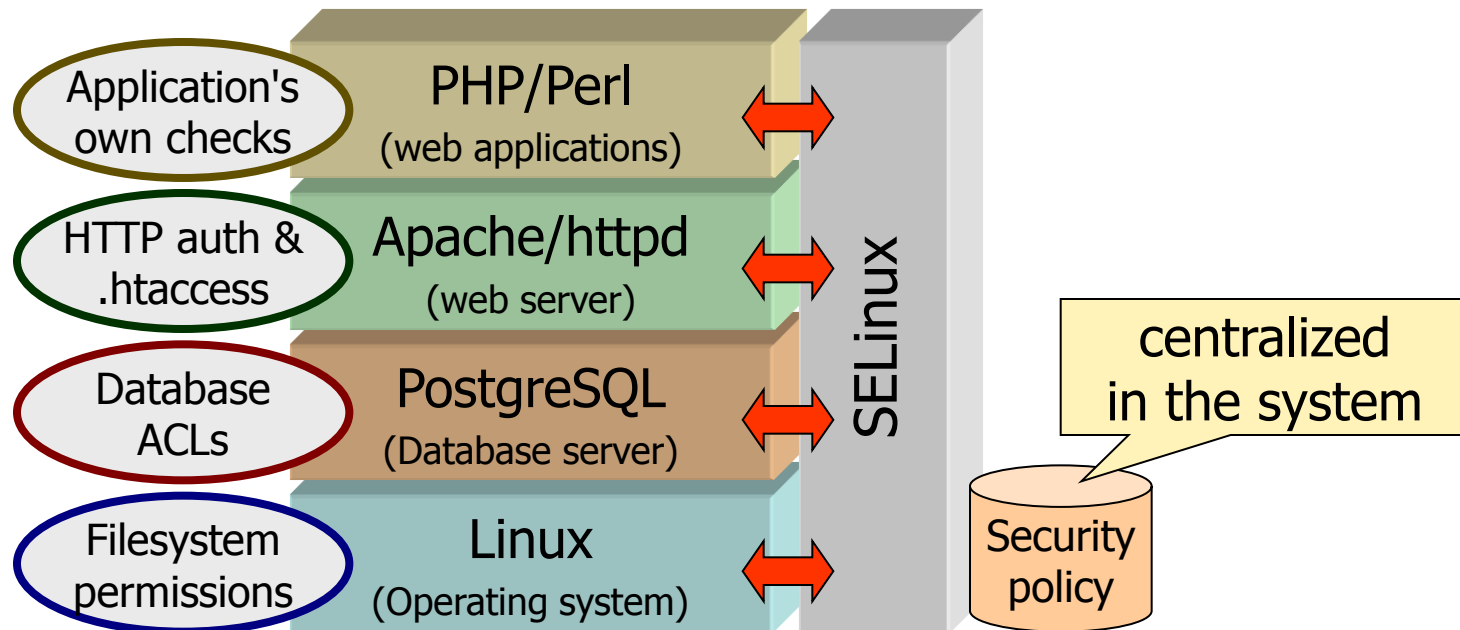
# Lack of conductor



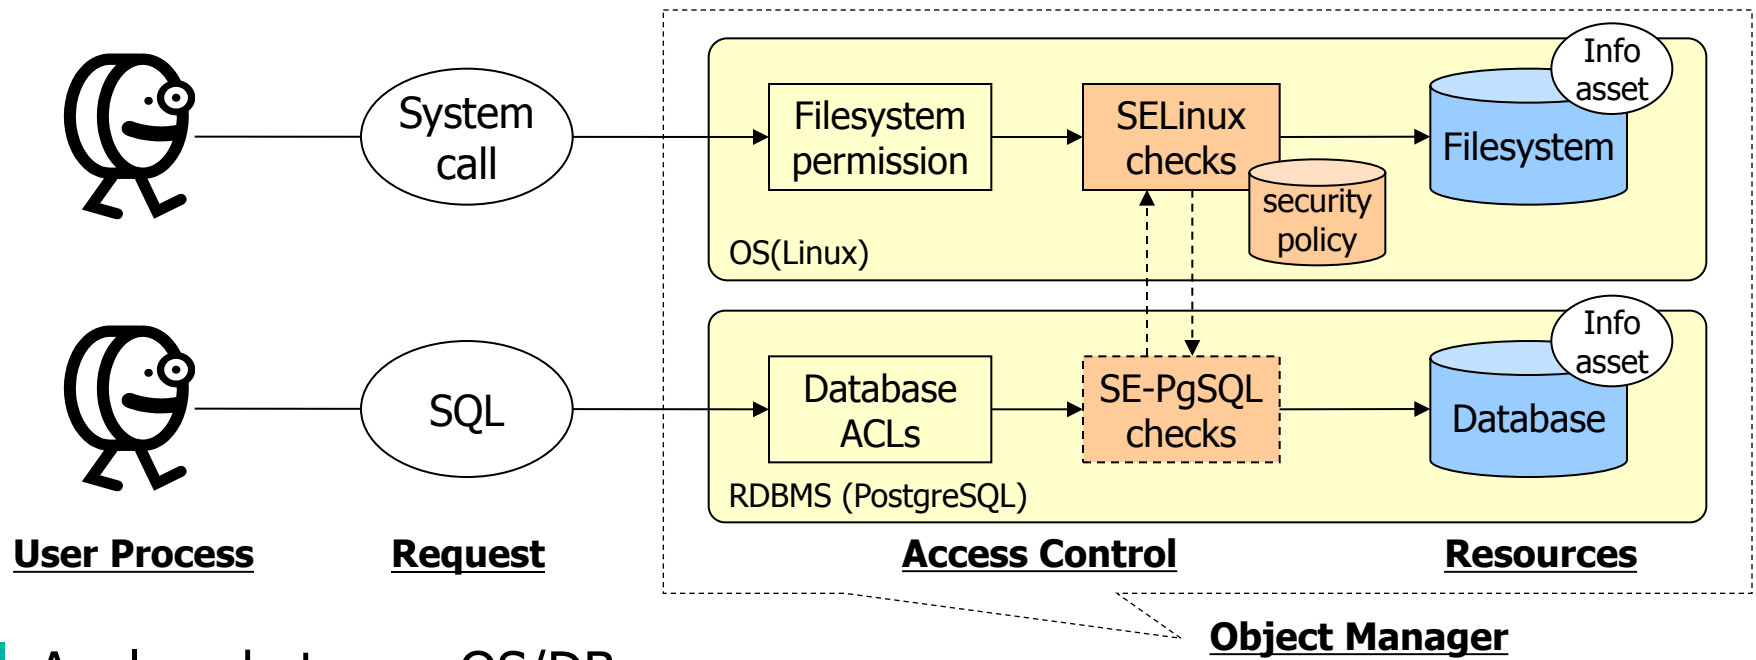
# LAPP/SELinux - concept

## SELinux performs as conductor

- System-wide privileges are assigned to all the users
- DB controls accesses based on the centralized policy
- ➡ It ensures **least-privilege** and **consistency** in access control.



# Perspective from the model (1/2)



## Analogy between OS/DB

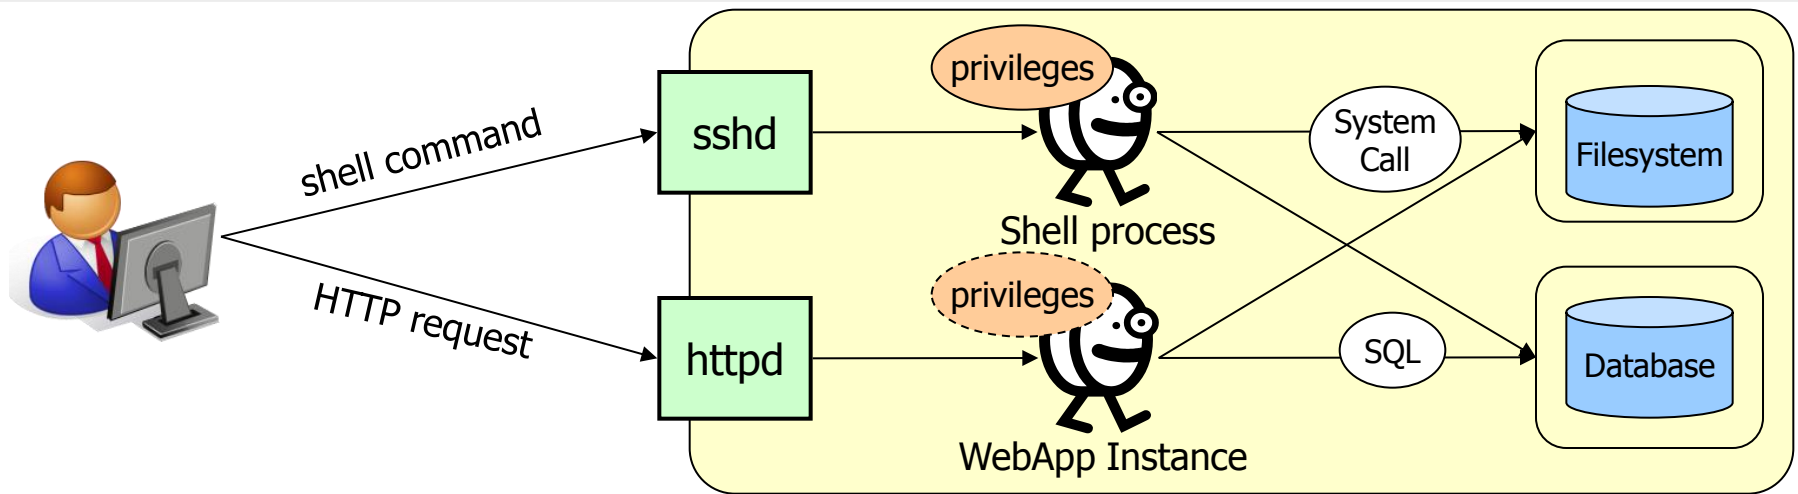
- Differences in the way to store and access information assets
- System-call for Filesystem, SQL for Database

## Role of access control

- It decides what is allowed or disallowed between users and resources, and controls the given requests based on the decision.
- ➡ Same basis (security policy) ensures system-wide consistency.



# Perspective from the model (2/2)



<u>User(Human)</u>	<u>Request(1st)</u>	<u>Authentication &amp; Authorization</u>	<u>User Agent</u>	<u>Request(2nd)</u>	<u>Resources</u>
--------------------	---------------------	---	-------------------	---------------------	------------------

## Analogy between shell and web

- User is a human; An user-agent performs instead of himself.
- User-agent must have correct privileges reflecting the actual human.

## Role of authentication & anthorization

- It identifies the human connected, and assigns their privileges.
  - sshd assigns user/group-id on the login shell before the execution.
  - Apache does not change privileges of the web-application instance.

# LAPP/SELinux - components

## SE-PostgreSQL

- A built-in enhancement of PostgreSQL
- Additional permission checks on the given queries according to the decision of SELinux
- It ensures consistency in access controls

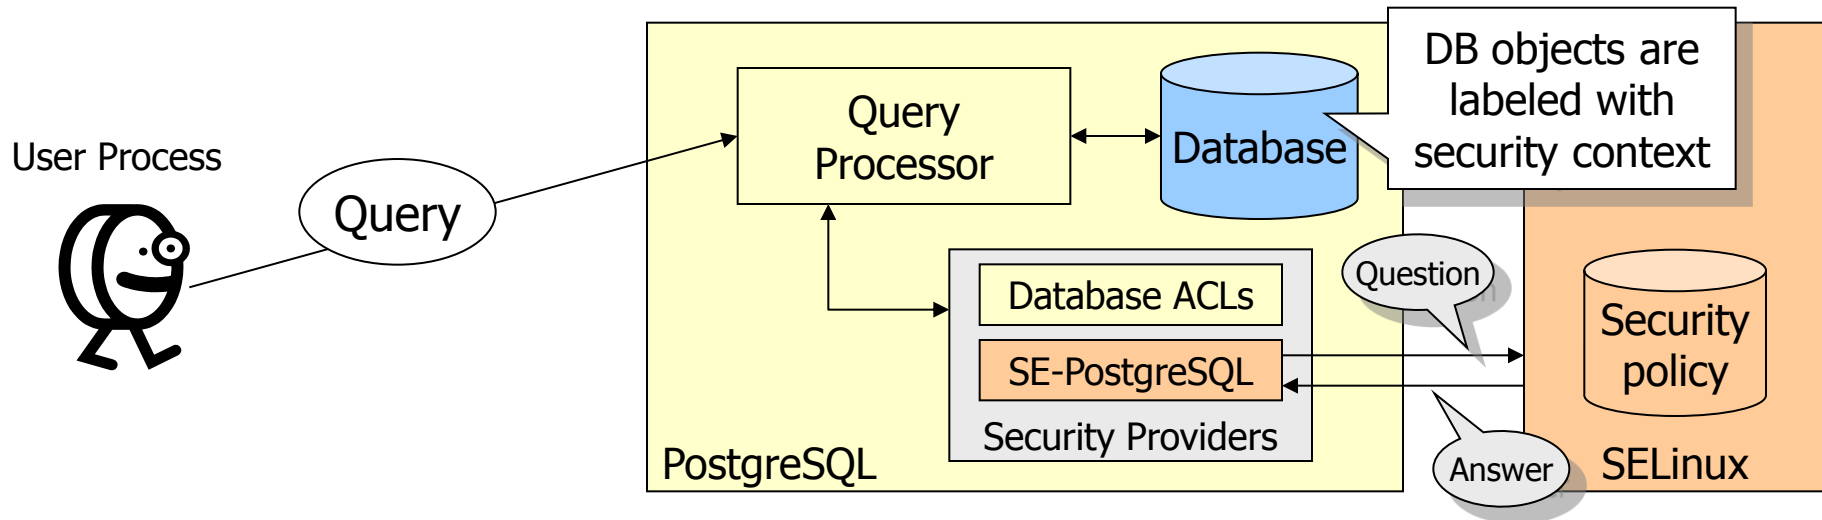
## Apache/SELinux Plus

- A loadable module of the Apache/httpd 2.2.x
- It assigns a security context of the contents handler based on http authentication.
- It ensures least-privilege in access control; with utilization of OS/DB

# Agenda

1. Background
2. SE-PostgreSQL
3. Apache/SELinux plus
4. Demonstration
5. Future Plans

# Architecture of SE-PostgreSQL



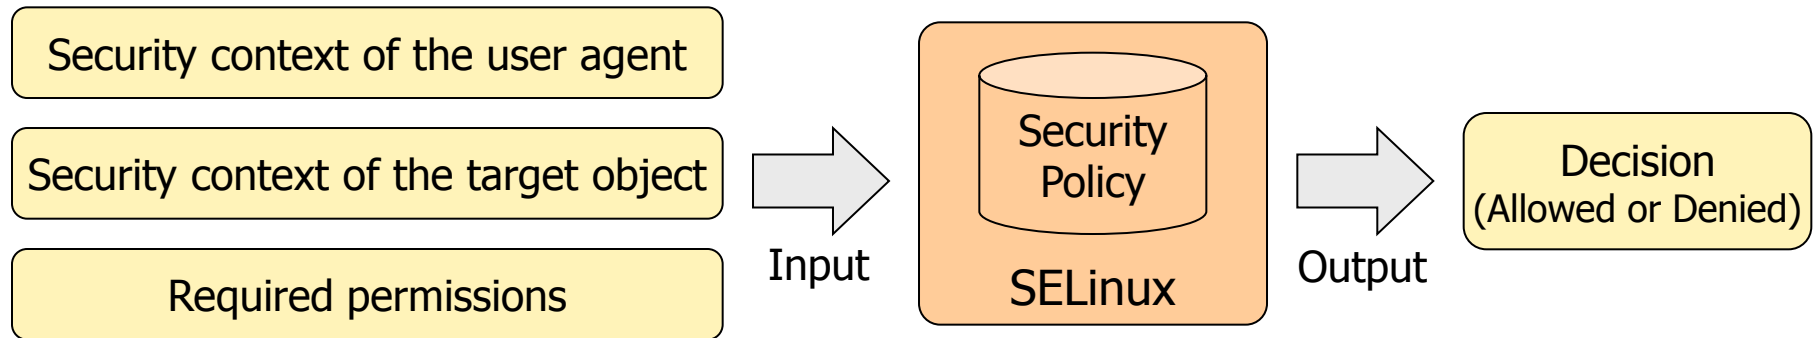
## Security Providers

- Common entrypoints of access control features; like database ACLs.
- SE-PostgreSQL shall be an optional security provider.

## SE-PostgreSQL

- It tells SELinux whether the given query is allowed to run;  
(Need to deliver a pair of security context of the client and objects)
- SELinux returns its decision,  
then SE-PostgreSQL raises an error if access violation.

# Decision making in access controls



## SELinux performs like a function

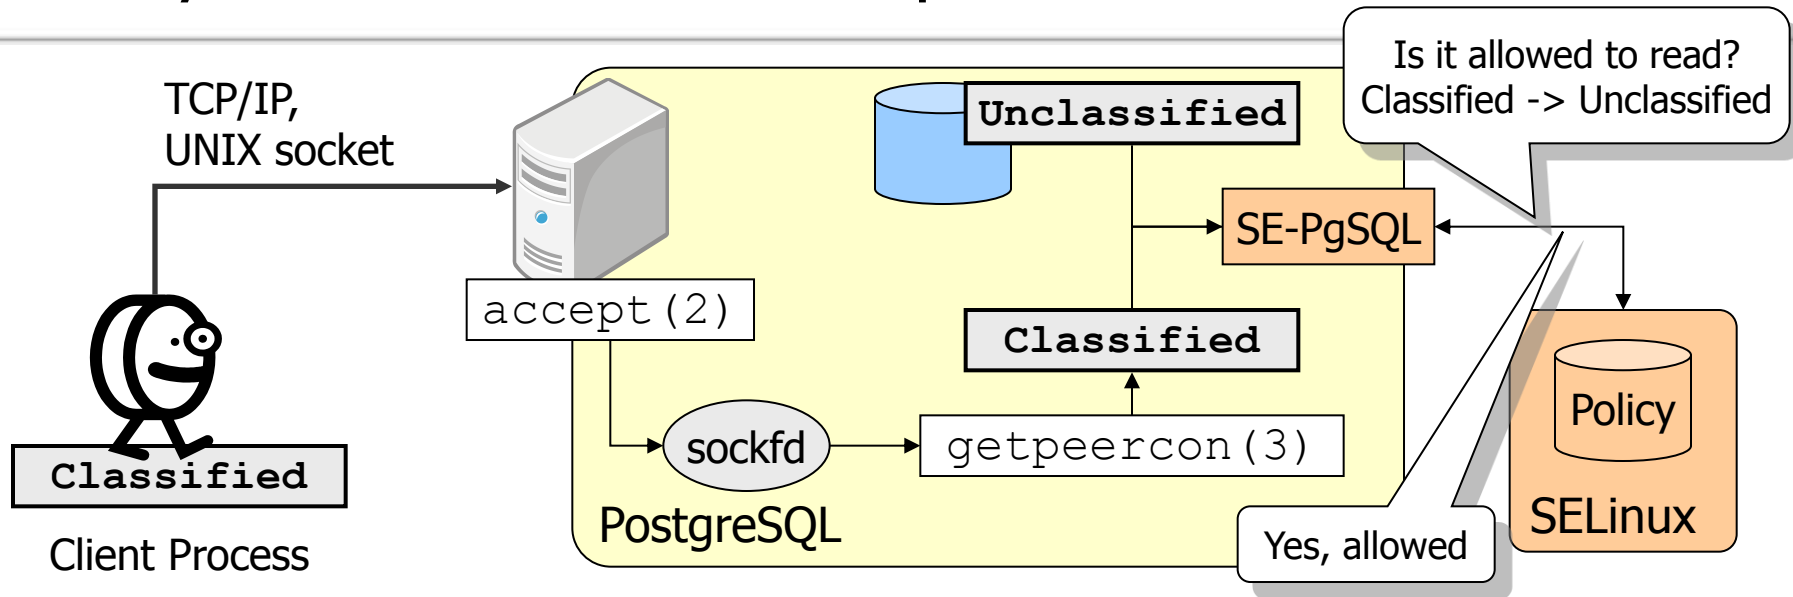
- It returns its decision for the given arguments.
- Kernel internally gives them to SELinux, and follows its decision.
- Userspace application can also utilize the mechanism, **as long as it can provide pair of the security context.**

## Security context

- A SELinux specified identifier of processes and any other objects.

```
Example)  system_u:system_r:httpd_t:s0  
          system_u:object_r:postgresql_db_t:s0
```

# Security context of the client process



## Labeled networks

- SELinux provides an API to obtain security context of the peer process.

```
int getpeercon(int sockfd, security_context_t *con);
```

- IPsec daemon exchanges the security context of peers prior to open the connection.
- Static fallback security context for non-SELinux'ed clients.
- It allows to identify the client process using security context.



# Security context of the database objects (1/2)

```
postgres=> SELECT security_label, * FROM drink;
          security_label          | id | name  | price
-----+-----+-----+-----
system_u:object_r:sepgsql_table_t:s0 | 1 | water | 110
system_u:object_r:sepgsql_table_t:s0 | 2 | tea   | 130
system_u:object_r:sepgsql_table_t:s0:c0 | 3 | coke  | 130
system_u:object_r:sepgsql_table_t:s0:c1 | 4 | coffee | 180
(4 rows)
```

## "security\_label" system column

- It represents the security context of tuples.
- The tuple of `pg_class` shows properties of table, so it means the security context of the table, for example.

## Default security context

- On insertion, the default one shall be assigned based on the policy.
- User can also provide an explicit one, instead of the default.

# Security context of the database objects (2/2)

## ALTER xxx SECURITY LABEL TO

- It allows to change security context of database objects.
- Use **UPDATE** statements for tuples within user tables.

```
postgres=> ALTER TABLE t SECURITY LABEL TO
            'user_u:object_r:sepgsql_ro_table_t:s0';
ALTER TABLE
```

## ALTER TABLE xxx SET WITH/WITHOUT SECURITY LABEL

- It allows to strip 'security\_label' system column, if not necessary.
- Reduce row-level control and storage consumption on the table.

```
postgres=> ALTER TABLE t SET WITHOUT SECURITY LABEL;
ALTER TABLE
postgres=> SELECT security_label, * FROM t;
ERROR:  column "security_label" does not exist
```

# Usage of SE-PostgreSQL (1/2)

```
postgres=# CREATE TABLE customer
          (id integer primary key, name text, credit text);
postgres=# ALTER TABLE customer ALTER credit SECURITY LABEL TO
          'system_u:object_r:sepgsql_secret_table_t:s0';
postgres=# INSERT INTO customer
          VALUES (1, 'kaigai', '1111-2222-3333-4444');
```

```
postgres=# SELECT * FROM customer;
LOG:  SELinux: denied { select } Ȳ
      scontext=staff_u:staff_r:staff_t:s0 Ȳ
      tcontext=system_u:object_r:sepgsql_secret_table_t:s0 Ȳ
      tclass=db_column name=customer.credit
```

**ERROR: SELinux: security policy violation**

```
postgres=# SELECT id, name FROM customer;
```

```
 id | name
----+-----
  1 | kaigai
(1 row)
```

Client was not allowed to select  
from the column labeled as  
sepgsql\_secret\_table\_t

# Usage of SE-PostgreSQL (2/2)

```
postgres=# SELECT security_label, * FROM;
```

security_label	id	name	price
system_u:object_r:sepgsql_table_t:Unclassified	1	water	100
system_u:object_r:sepgsql_table_t:Classified	2	coke	120
system_u:object_r:sepgsql_ro_table_t:Classified	3	juice	140
system_u:object_r:sepgsql_ro_table_t:Unclassified	4	coffee	180
staff_u:object_r:sepgsql_table_t:Unclassified	5	beer	240

## On SELECT

- All the tuples are visible for Classified user, but Classified tuples are not visible Unclassified user.


## On UPDATE/DELETE

- Also, Classified tuples are updatable/deletable by Classified users.
- And, Read-only tuples are not updatable by confined users.

## On INSERT

- A default security context shall be assigned on the new tuple, and checks privilege to insert it.

# System-wide consistency in access controls




SELinux provides its access control decision for ANY subsystems

- ✓ Linux kernel enforces the decision on accesses to filesystem objects, and etc...
- ✓ SE-PostgreSQL enforces the decision on accesses to database objects.

➡ Eventually, the centralized security policy controls all the accesses

# Performance - SE-PostgreSQL



2~4% of trade-off in performance

- userspace AVC minimizes the number of kernel invocations

Environments


- CPU Xeon (2.33GHz) Dual, Mem: 2GB (shared\_buffer=512m)
- measured by `pgbench -c 2 -t 200000`



# Agenda

1. Background
2. SE-PostgreSQL
3. Apache/SELinux plus
4. Demonstration
5. Future Plans

# Who's privileges should be checked?



## Authentication, but no authorization

- Apache can check client's Web-ID/PASS (BASIC or DIGEST).
- 403 Error, or Apache launches web-application handlers.

## Problem

- Web-application performs with identical privilege of daemon process.
- It means OS/RDBMS cannot distinguish individual web-users.
- Web-applications have to work always correctly?

It means web-applications have to be bugs/vulnerabilities free? ☹

# Web users

The screenshot shows a web browser window with the URL `https://www.amazon.com/gp/yourstore?ie=UTF8&ref=pd_irl_gw&signIn=1`. The main page is the Amazon.com sign-in interface. On the left, there is a sidebar with a yellow box asking "Don't have a Yahoo! ID? Signing up is easy. Sign up for Yahoo!". Below this, another yellow box asks "Already have a Yahoo! ID? Sign in." and "Are you protected? Create your sign-in seal. (Why?)". The sign-in form includes fields for "Yahoo! ID:" (containing `var@example.com`) and "Password:". There are checkboxes for "Keep me signed in" and a "Sign In" button. At the bottom of the sidebar, it says "I can't access my account | Help".

The main content area of the browser shows the Amazon.com logo and navigation links. Below the navigation bar, there is a "Sign In" section with the text "What is your e-mail address?". The email field contains `foo@example.com`. Below this, it asks "Do you have an Amazon.com password?". There are two radio buttons: "No, I am a new customer." and "Yes, I have a password:". The "Yes" option is selected, and the password field contains `*****`. There is a checkbox for "パスワードを記憶する(R)" (Remember password).

Overlaid on the right side of the browser window is a "Google アカウント" (Google Account) sign-in form. It has fields for "メール:" (containing `baz@example.com`) and "パスワード:" (containing `*****`). There is a checkbox for "ログイン状態を保持する" (Keep me signed in) and a "ログイン" (Login) button. Below the form, it says "アカウントにアクセスできない場合" (If you cannot access your account).

Overlaid on the bottom center of the browser window is a small dialog box titled "saba.linux.bs1.fc.nec.co.jp に接続" (Connect to saba.linux.bs1.fc.nec.co.jp). It asks "Please enter your ID and Password". The "ユーザー名(U):" (Username) field contains `hoge` and the "パスワード(P):" (Password) field contains `*****`. There is a checkbox for "パスワードを記憶する(R)" (Remember password). At the bottom of the dialog are "OK" and "キャンセル" (Cancel) buttons.

# Not web-users

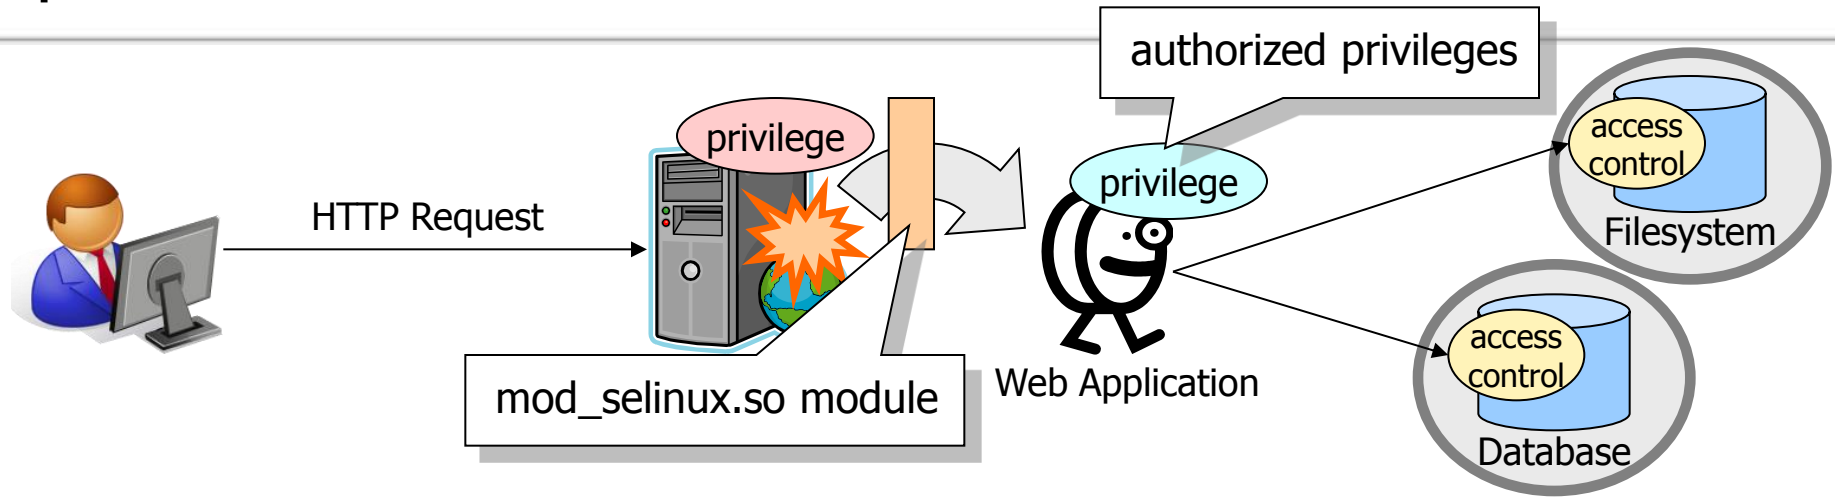
```
[kaigai@saba ~]$ ps -C httpd -o label,pid,user,group,comm
```

LABEL	PID	USER	GROUP	COMMAND
system_u:system_r:httpd_t:s0	25132	root	root	httpd
system_u:system_r:httpd_t:s0	25136	apache	apache	httpd
system_u:system_r:httpd_t:s0	25137	apache	apache	httpd
system_u:system_r:httpd_t:s0	25138	apache	apache	httpd
system_u:system_r:httpd_t:s0	25139	apache	apache	httpd
system_u:system_r:httpd_t:s0	25140	apache	apache	httpd
system_u:system_r:httpd_t:s0	25141	apache	apache	httpd
system_u:system_r:httpd_t:s0	25142	apache	apache	httpd
system_u:system_r:httpd_t:s0	25143	apache	apache	httpd
system_u:system_r:httpd_t:s0	25144	apache	apache	httpd

Security context of the httpd daemon  
used to access controls in SELinux

UNIX Uid/Gid of the httpd daemon  
used to discretionary access controls

# Apache/SELinux Plus



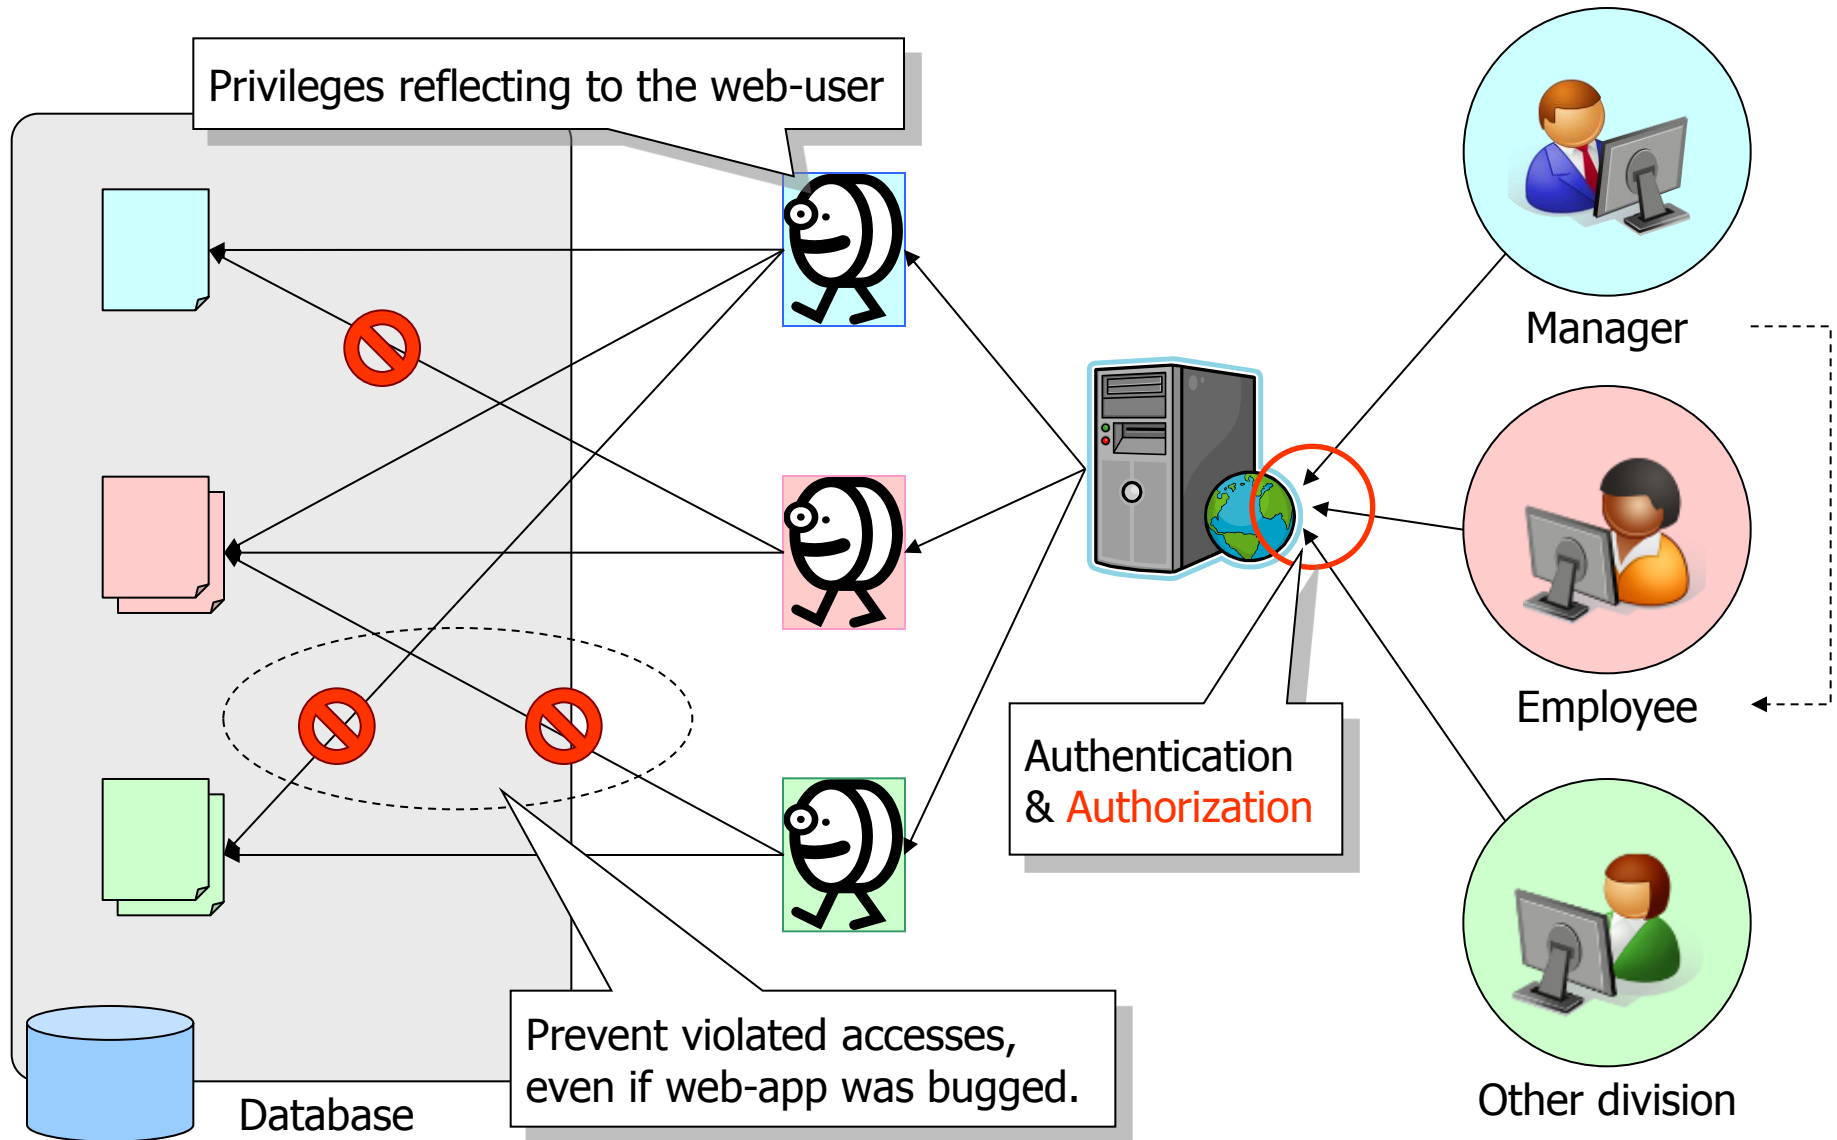
## Terms

- Authentication is a function of identifying the connected user.
- Authorization is a function of assigning the rights to resources.

## Apache/SELinux Plus (mod\_selinux.so)

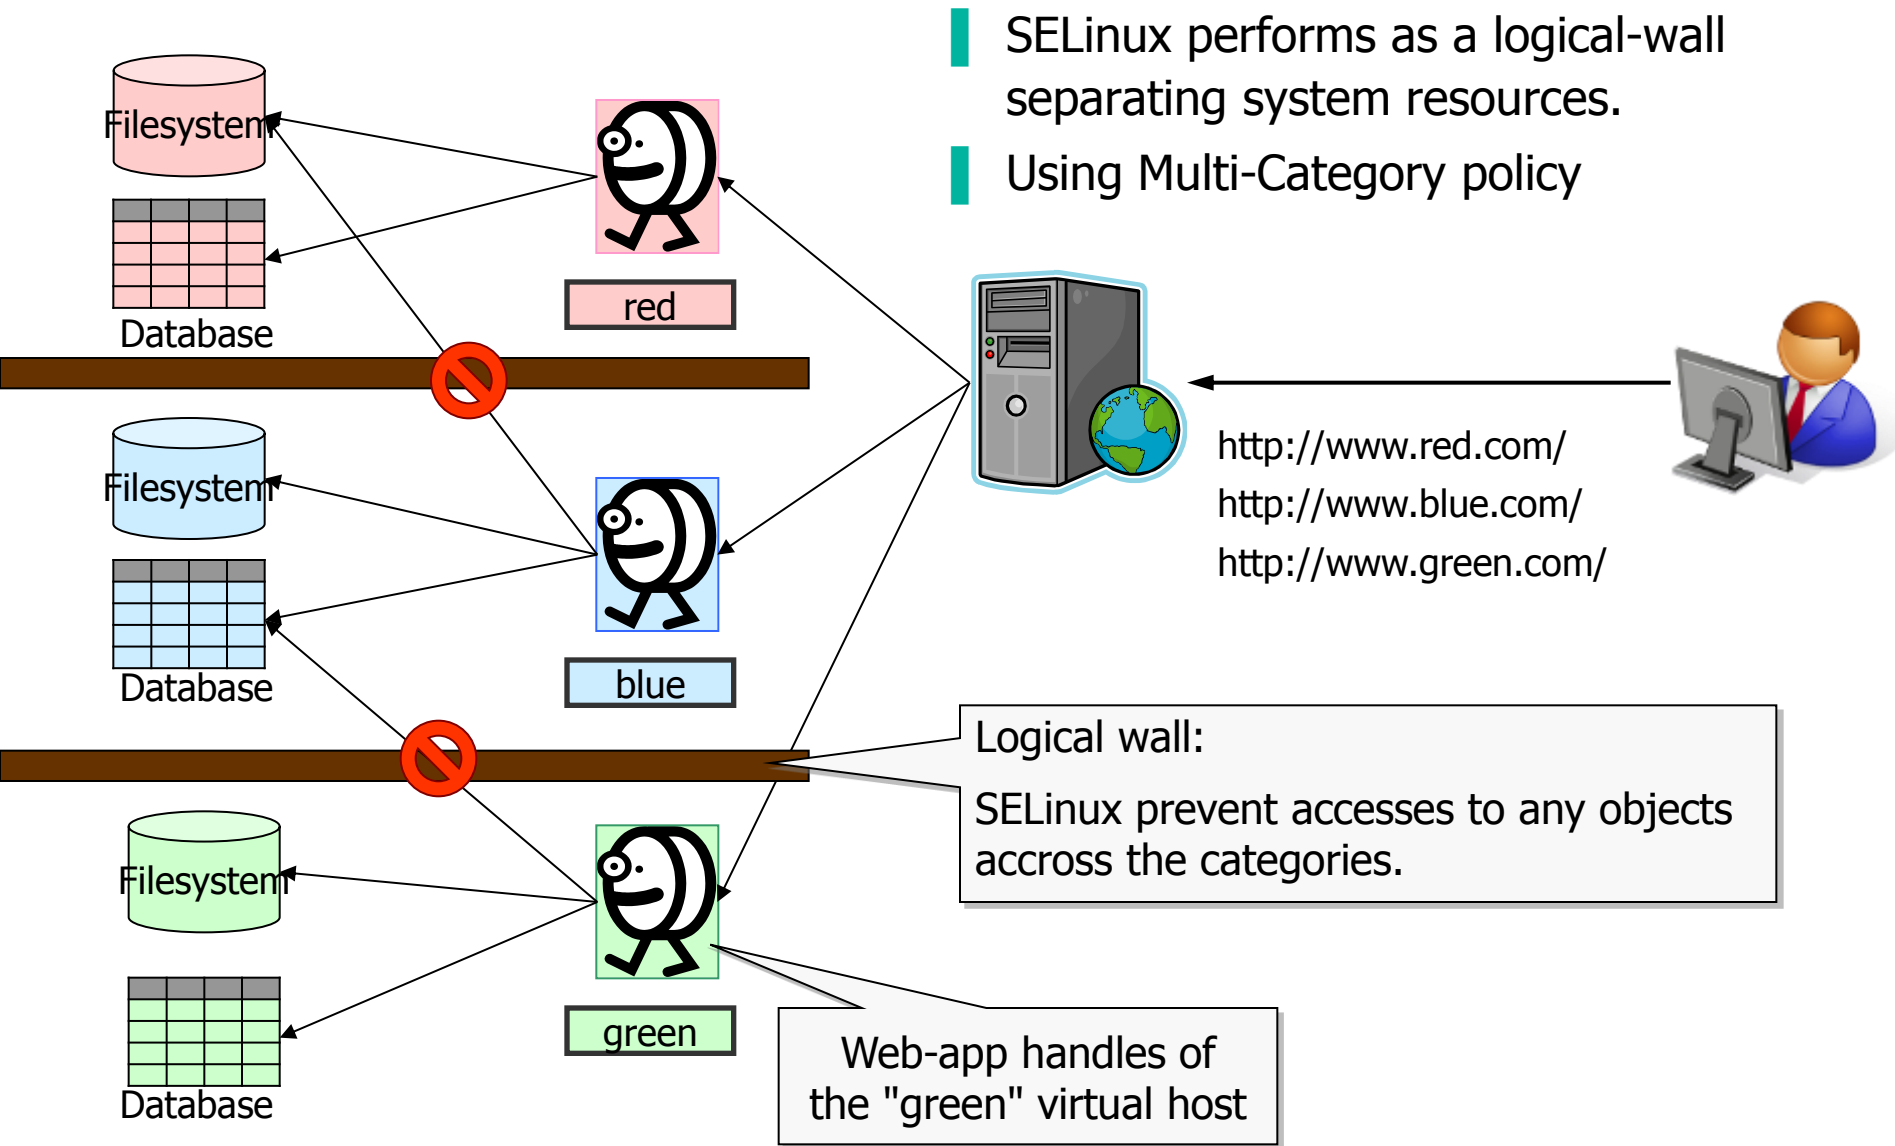
- It assigns a corresponding security context based on HTTP authentication prior to web-application launches.
- It enables to confine web-application's accesses.
- Unlike UNIX, no root capabilities are needed to change privileges.

# System-image (1/2) : Per web-user separation






# System image (2/2) : Per virtual host separation



# Performance - Apache/SELinux Plus



The cost to assign privileges is relatively large in lightweight request.

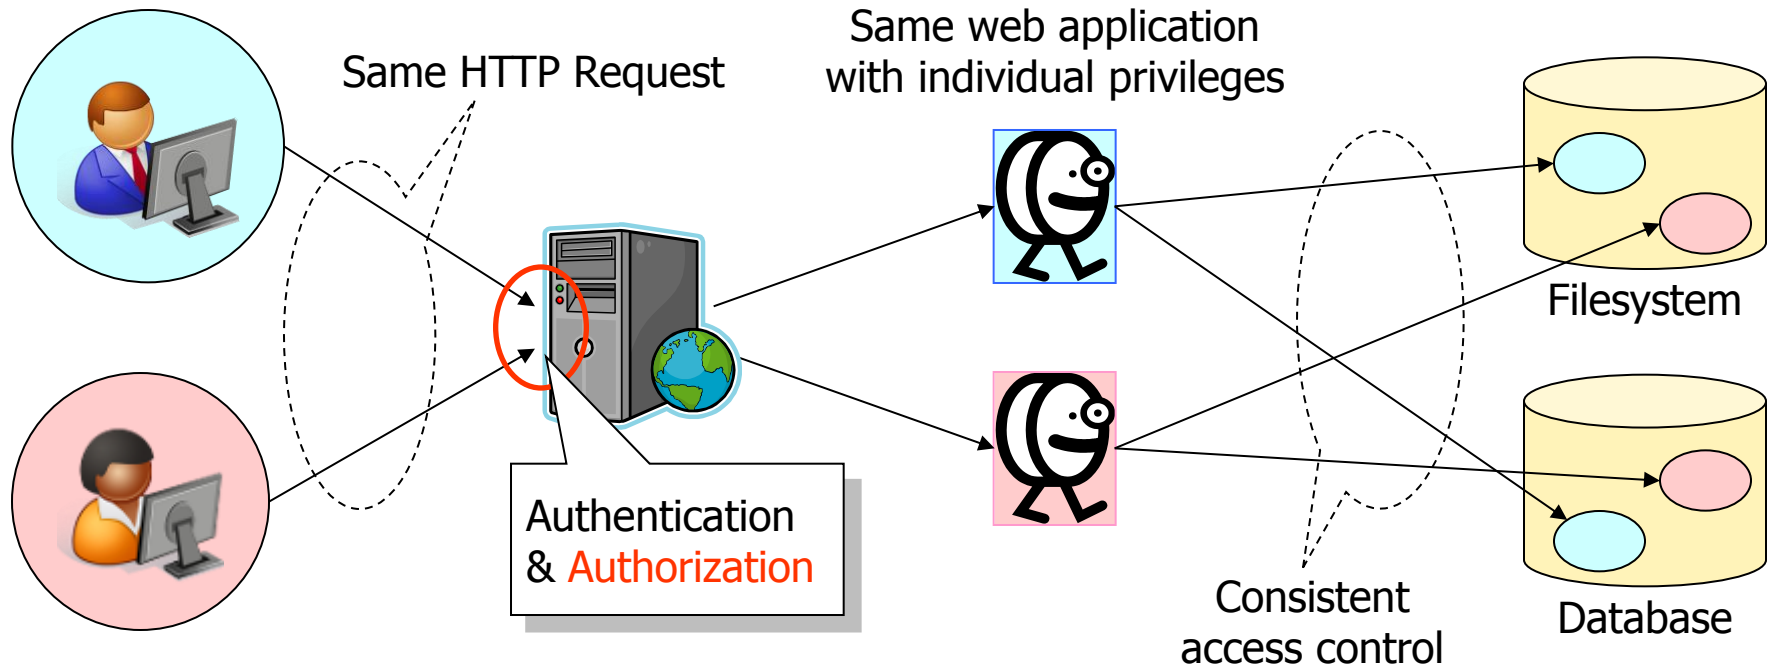
Less differences in our major target (Web+DB applications)

➡ Other steps obscures the cost to assign privileges.

# Agenda

1. Background
2. SE-PostgreSQL
3. Apache/SELinux plus
4. **Demonstration**
5. Future Plans

# Demonstration



Apache/SELinux Plus launches a PHP script with individual privileges.

The PHP script can access both of filesystem and database.

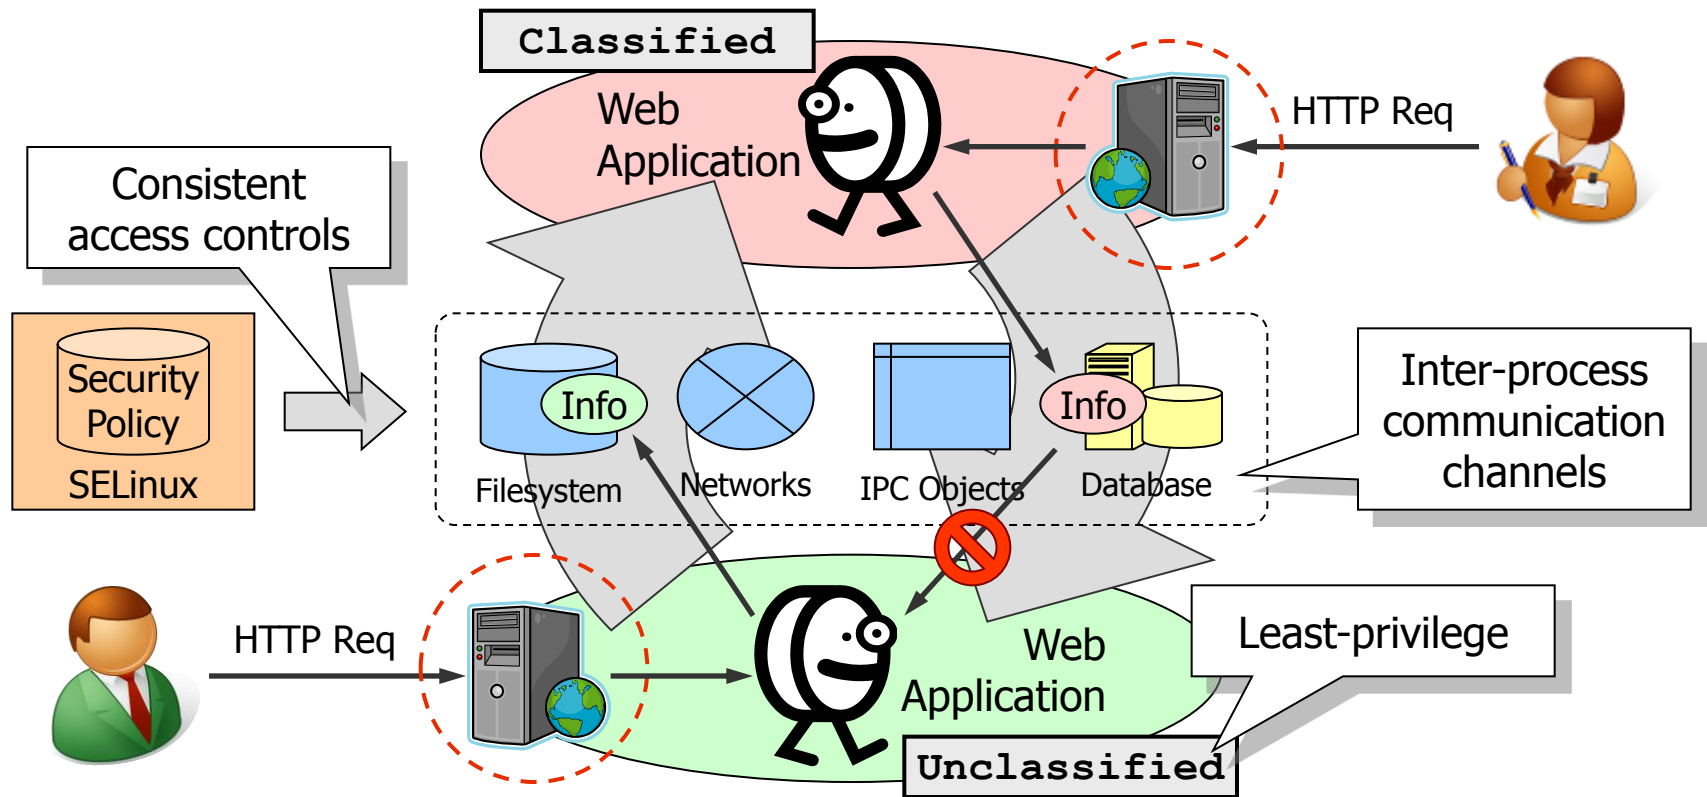
- Linux applies access controls on filesystems
- PostgreSQL applies access controls on databases

➡ Consistent access controls, although different mechanisms decide it.

# Agenda

1. Background
2. SE-PostgreSQL
3. Apache/SELinux plus
4. Demonstration
5. Future Plans

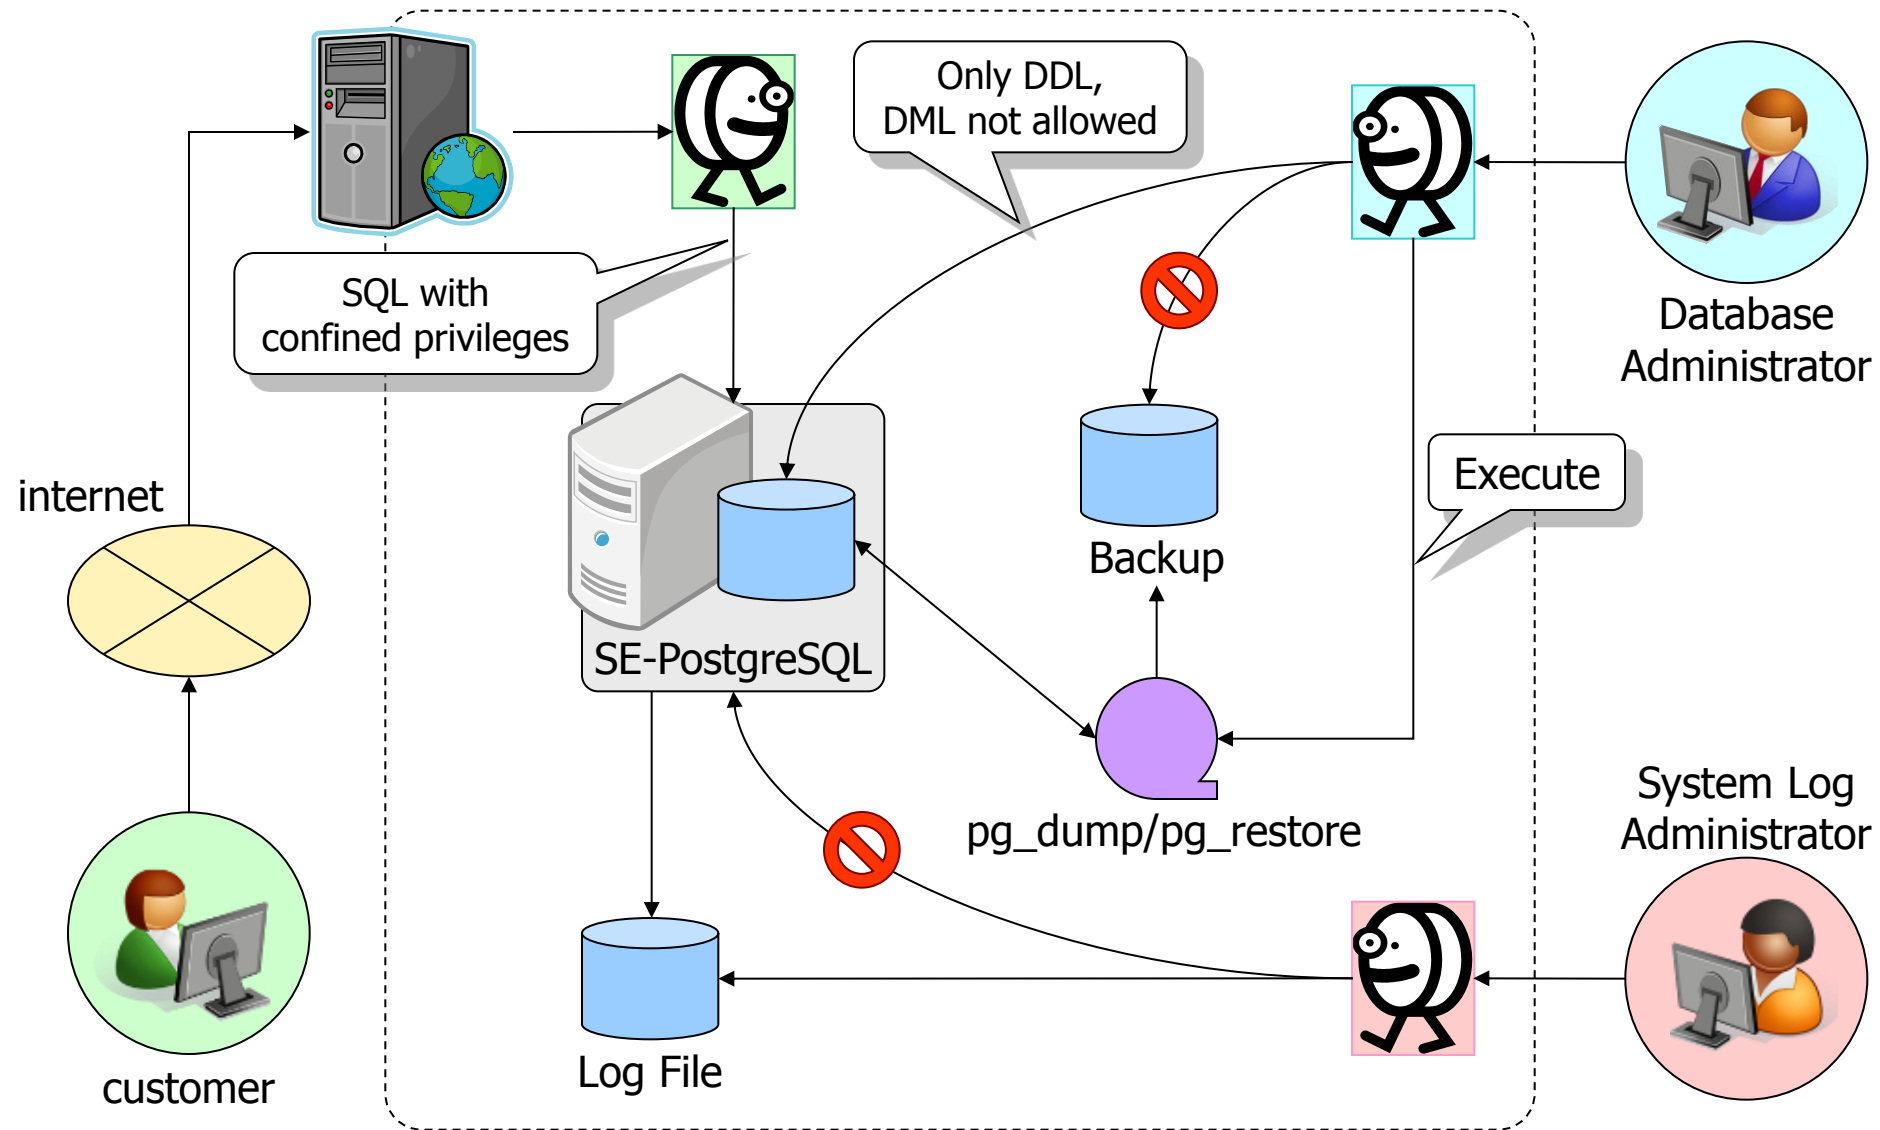
# Conceptual diagram of LAPP/SELinux



- SE-PostgreSQL ensures system-wide **consistency** in access controls.
- Apache/SELinux Plus ensures **least-privilege** on web-applications.
- ➡ LAPP/SELinux provides a secure web-application platform.

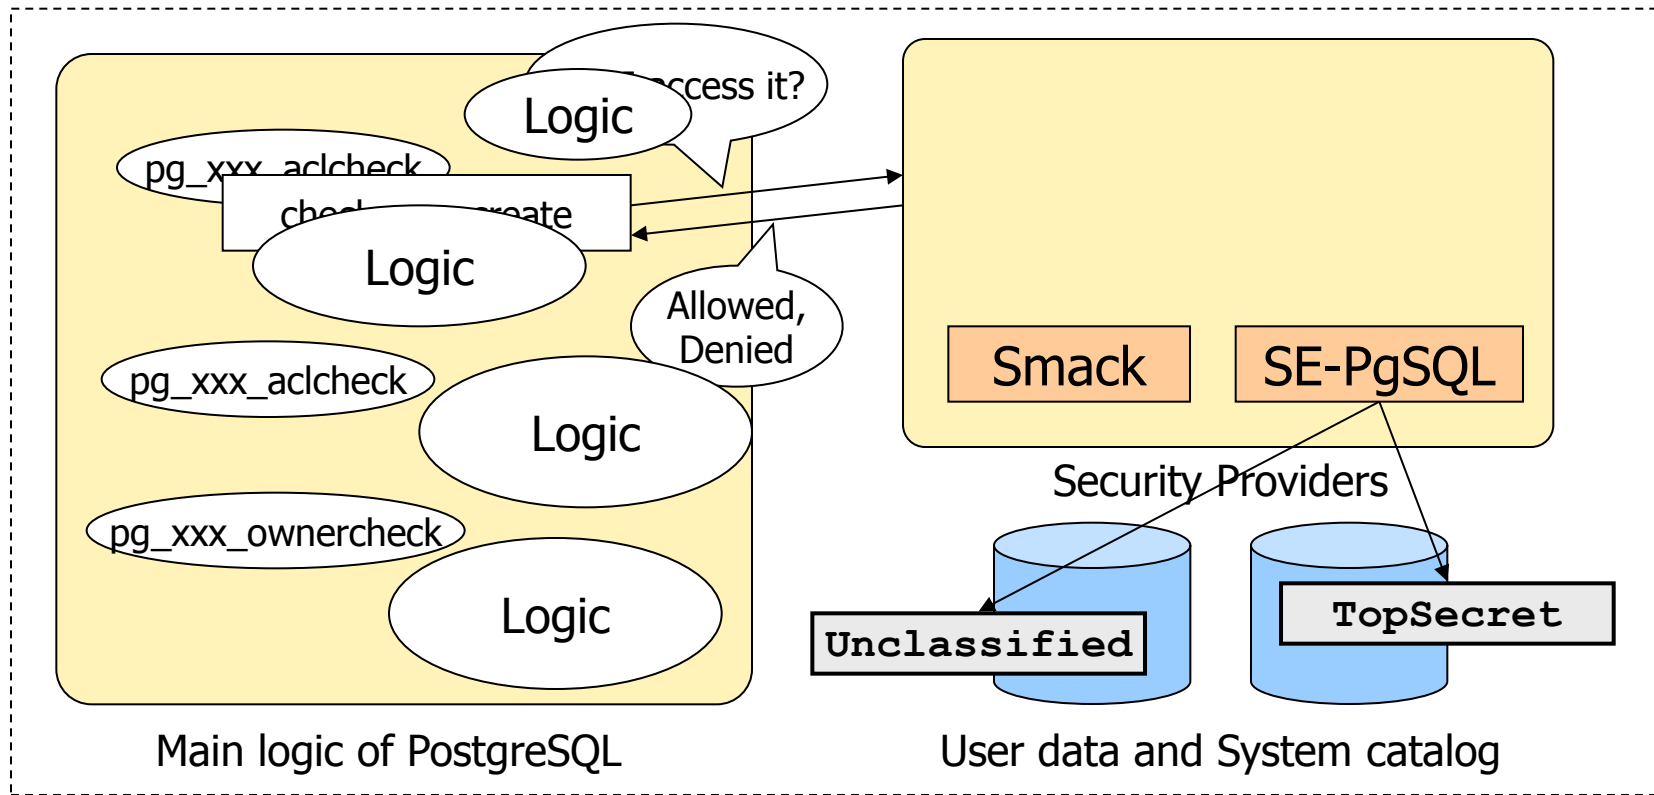


# Idea: Role Based Access Control System



# Plan to upstream: SE-PostgreSQL

- Access control reworks
- Add security label support
- Add an optional security provider



# Summary of LAPP/SELinux

## Background

- Web Application's security is Hot issue now.

## Key concept

- Utilize SELinux as conductor of access control

## Key components

- SE-PostgreSQL
- Apache/SELinux Plus

## Road To **SE-PostgreSQL** being Upstreamed

- External Security Providers
- Security Label Support
- SELinux support; as one of the security providers

➡ Here we go! Let's join us on v9.1 development!

# Any Questions?



# Thank you!



Empowered by Innovation

**NEC**

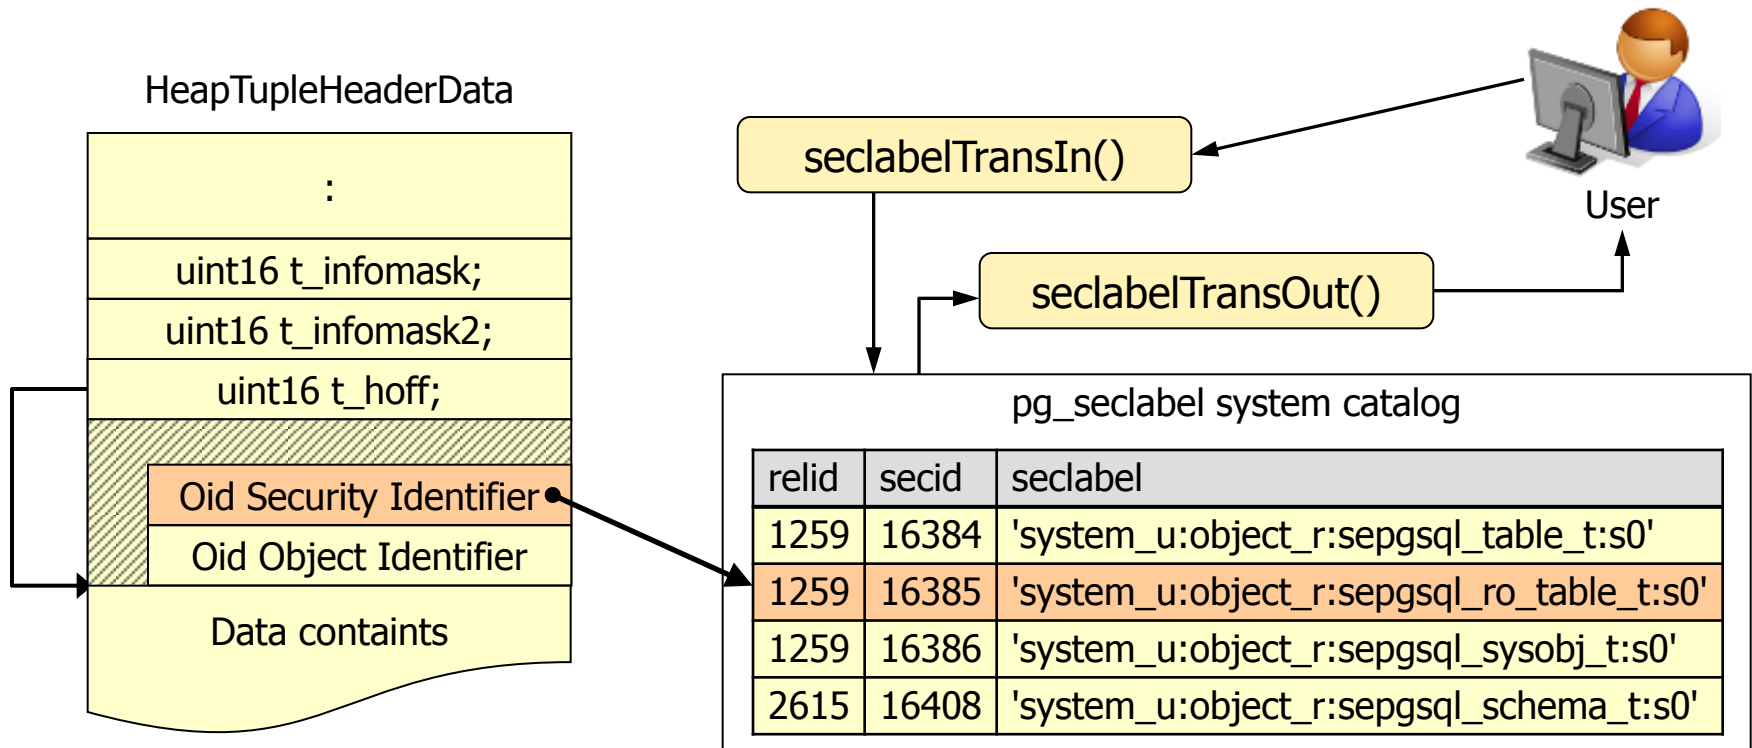
# Appendix



# Management of the security labels

## Data format

- A tuple has its security context as an object identifier (4-bytes).
  - It minimizes the waste of storage to store security context.
  - It allows to lookup avc cached without text comparison.
- pg\_seclabel system catalog holds its text representation.



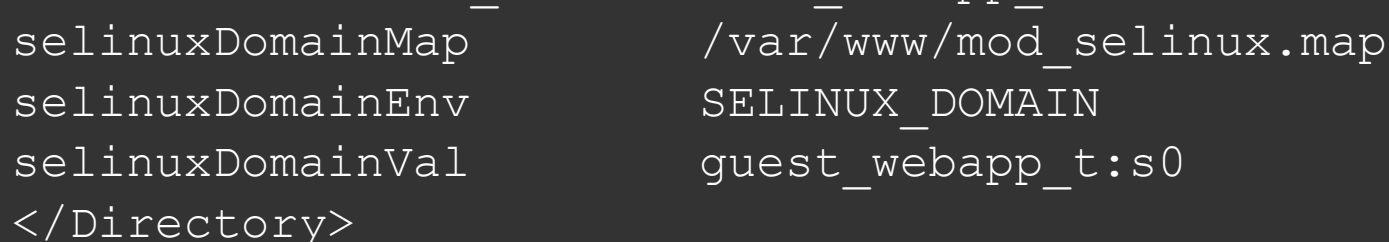


# Apache/SELinux Plus configuration (1/2)

```
# Apache/SELinux Plus configuration
# -----
LoadModule selinux_module modules/mod_selinux.so

selinuxServerDomain      *:s0

<Directory "/var/www/html">
SetEnvIf Remote_Addr "192.168.1.[0-9]+$"  ¥
    SELINUX_DOMAIN=user_webapp_t:s0
selinuxDomainMap          /var/www/mod_selinux.map
selinuxDomainEnv          SELINUX_DOMAIN
selinuxDomainVal          guest_webapp_t:s0
</Directory>
```



Order to be applied

A pair of the http authorized  
username and security context

```
# Apache/SELinux Plus user-mapping
# -----
foo      user_webapp_t:s0:c0
var      user_webapp_t:s0:c1
baz      user_webapp_t:s0:c2
```

# Apache/SELinux Plus configuration (2/2)

```
# Apache/SELinux Plus (Per VirtualHost Separation)
# -----
LoadModule selinux_module modules/mod_selinux.so

selinuxServerDomain      *:s0-s0:c0.c1

<VirtualHost *:80>
DocumentRoot             /var/www/html
ServerName               red.example.com
selinuxDomainVal         *:s0:c0
</VirtualHost>
```

Web-server process MUST dominate all the categories.

```
<VirtualHost *:80>
DocumentRoot             /var/www/html
ServerName               blue.example.com
selinuxDomainVal         *:s0:c1
</VirtualHost>
```

It assigns c1 category for all the HTTP requests including anonymous ones.