



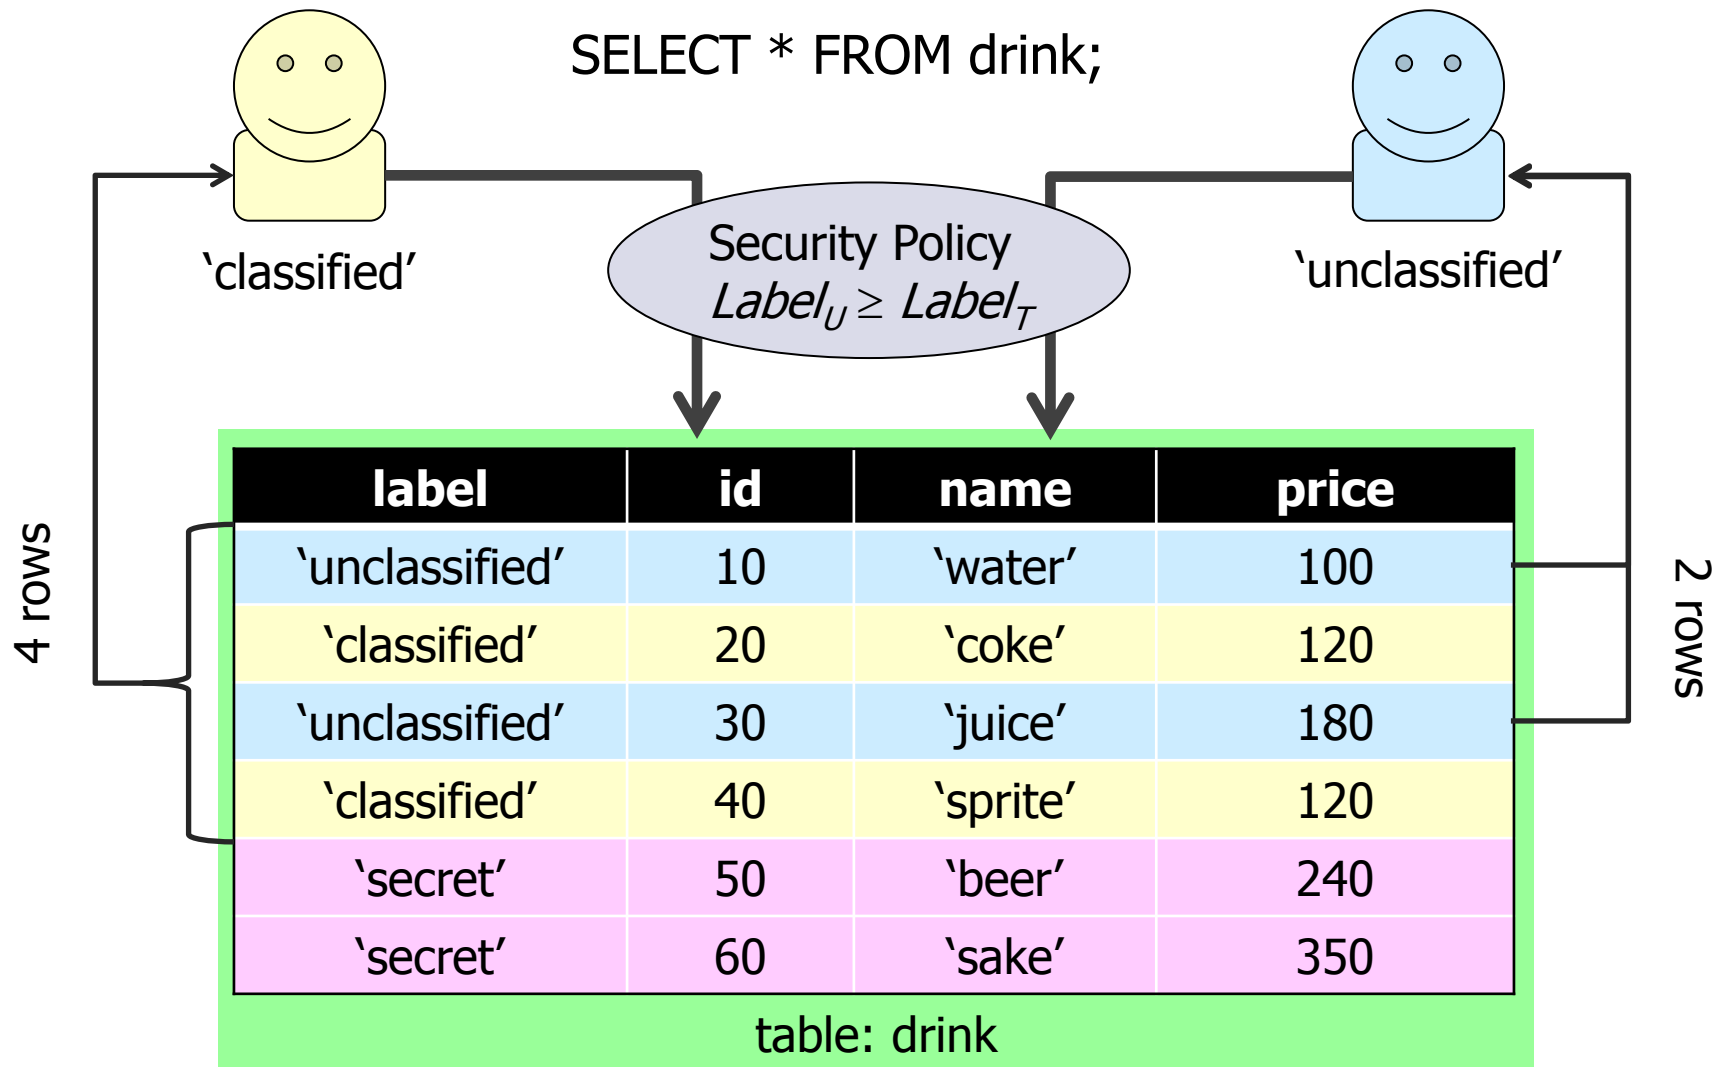
# Row Level Security

KaiGai Kohei <kaigai@kaigai.gr.jp>

tw: @kkaigai



# How RLS should work (1/2)





# How RLS should work (2/2)



shop\_id = 100

SELECT \* FROM drink NATURAL JOIN drink\_order

id	name	price	shop_id	quantum	data
10	'water'	100	100	8	2013-02-16
30	'juice'	180	100	10	2013-02-18

id	name	price
10	'water'	100
20	'coke'	120
30	'juice'	180
40	'sprite'	120
50	'beer'	240
60	'sake'	350

table: drink

Security Policy  
 $shop\_id_U = shop\_id_T$

id	shop_id	quantum	date
10	100	8	2013-02-16
20	200	5	2013-02-17
10	200	6	2013-02-18
30	100	10	2013-02-18

table: drink\_order



# WHERE is simple solution?

```
postgres=> CREATE VIEW soft_drink AS
           SELECT * FROM drink WHERE price < 200;
CREATE VIEW
postgres=> GRANT SELECT ON soft_drink TO public;
GRANT
postgres=> SET SESSION AUTHORIZATION bob;
SET
postgres=> SELECT * FROM soft_drink;
 id |  name  | price
----+-----+-----
 10 | water  |   100
 20 | coke   |   120
 30 | juice  |   180
 40 | sprite |   120
(4 rows)

postgres=> SELECT * FROM drink;
ERROR:  permission denied for relation drink
```



# Nightmare of Leaky View (1/3)

```
postgres=> SELECT * FROM soft_drink WHERE f_leak(name);
```

```
NOTICE:  f_leak => water
```

```
NOTICE:  f_leak => coke
```

```
NOTICE:  f_leak => juice
```

```
NOTICE:  f_leak => sprite
```

```
NOTICE:  f_leak => beer
```

```
NOTICE:  f_leak => sake
```

id	name	price
10	water	100
20	coke	120
30	juice	180
40	sprite	120

(4 rows)



## Nightmare of Leaky View (2/3)

```
postgres=> CREATE OR REPLACE FUNCTION f_leak (text)
  RETURNS bool COST 0.000001 AS
  $$
  BEGIN
  RAISE NOTICE 'f_leak => %', $1;
  RETURN true;
  END
  $$ LANGUAGE plpgsql;
CREATE FUNCTION
```

```
postgres=> EXPLAIN(costs off)
  SELECT * FROM soft_drink WHERE f_leak(name);
      QUERY PLAN
```

```
-----
Seq Scan on drink
  Filter: (f_leak(name) AND (price < 200))
(2 rows)
```

f\_leak() のコストは '<' 演算子より小さい



# Nightmare of Leaky View (3/3)

```
postgres=> CREATE VIEW v_both AS
  SELECT * FROM t_left JOIN t_right ON a = x
  WHERE b like '%hoge%';
CREATE VIEW
```

```
postgres=> EXPLAIN (COSTS OFF)
  SELECT * FROM v_both WHERE f_leak(y);
QUERY PLAN
```

-----

Hash Join

Hash Cond: (t\_left.x = t\_right.a)

-> Seq Scan on t\_left

**Filter: f\_leak(y)**

-> Hash

-> Seq Scan on t\_right

Filter: (b ~~ '%hoge%'::text)

(7 rows)

f\_leak()の実行は、  
t\_left だけに依存



# Problem to be tackled

```
SELECT * FROM v_both WHERE f_leak(y);
```




```
SELECT * FROM (  
    SELECT * FROM t_left JOIN t_right ON a = x  
    WHERE b like '%hoge%'  
) AS v_both WHERE f_leak(y)
```

- 問題の本質は、クエリ最適化が条件句を評価する順序を入れ替えてしまう事。
- 少なくともセキュリティ目的のビューなら、サブクエリの境界を跨いで条件句を移動させてはならない。





# Security Barrier (1/3)





## Security Barrier (2/3)

```
postgres=> CREATE OR REPLACE VIEW soft_drink
           WITH (security_barrier)
           AS SELECT * FROM drink WHERE price < 200;

CREATE VIEW
postgres=> SET SESSION AUTHORIZATION bob;
SET
postgres=> SELECT * FROM soft_drink WHERE f_leak(name);
NOTICE:  f_leak => water
NOTICE:  f_leak => coke
NOTICE:  f_leak => juice
NOTICE:  f_leak => sprite
 id | name  | price
----+-----+-----
 10 | water |   100
 20 | coke  |   120
 30 | juice |   180
 40 | sprite |   120
(4 rows)
```



# Security Barrier (3/3)

```
postgres=> EXPLAIN (costs off)
            SELECT * FROM soft_drink WHERE f_leak(name);
            QUERY PLAN
-----
Subquery Scan on soft_drink
  Filter: f_leak(soft_drink.name)
    -> Seq Scan on drink
        Filter: (price < 200)
(4 rows)
```

- CREATE VIEW ... WITH (security\_barrier) AS ...
- security\_barrier属性を持つVIEWの内側へは、条件句を push-down させない。
  - 利用者が意図した通りの順序で条件句が評価される
  - 代償として、最適な実行性能は得られなくなる



# Security-performance trade off

```
postgres=> CREATE VIEW my_team WITH (security_barrier)
           AS SELECT * FROM employee WHERE boss = current_user;
CREATE VIEW
postgres=> EXPLAIN (costs off)
           SELECT * FROM my_team WHERE id = 100;
           QUERY PLAN
-----
Subquery Scan on my_team
  Filter: (my_team.id = 100)
    -> Seq Scan on employee
         Filter: (boss = "current_user"())
(4 rows)
```

- **id = 100 を使ってインデックススキャンができるはず**
- **だが、security\_barrier属性のために、先にemployeeを全件スキャン → その後で id = 100 を評価**



# Leakproof Function (1/3)





## Leakproof Function (2/3)

- 副作用の無い (= 間違いなく安全な) 関数なら、  
security\_barrierビューの内側に押し込んでも大丈夫  
→ それを示すのが、LEAKPROOF 属性

```
postgres=# CREATE FUNCTION nabeatsu(integer)
           RETURNS bool LEAKPROOF AS
$$
BEGIN
    IF ($1 % 3 = 0) THEN RETURN true; END IF;
    WHILE $1 > 0 LOOP
        IF ($1 % 10 = 3) THEN RETURN true; END IF;
        $1 = $1 / 10;
    END LOOP;
    RETURN false;
END
$$ LANGUAGE plpgsql;
CREATE FUNCTION
```



# Leakproof Function (3/3)

```
postgres=> EXPLAIN (costs off)
            SELECT * FROM my_team WHERE nabeatsu(id);
            QUERY PLAN
-----
Seq Scan on employee
  Filter: ((boss = "current_user"()) AND nabeatsu(id))
(2 rows)
```

- デフォルトLEAKPROOFなビルトイン関数もある
  - integer対integer の 等価演算子など

```
postgres=> EXPLAIN (costs off)
            SELECT * FROM my_team WHERE id = 300;
            QUERY PLAN
-----
Index Scan using employee_pkey on employee
  Index Cond: (id = 300)
  Filter: (boss = "current_user"())
(3 rows)
```



# In case of Oracle

-----					
Id		Operation	Name	Rows	Bytes
-----					
0		SELECT STATEMENT		3	81
* 1		VIEW	V	3	81
* 2		HASH JOIN		3	120
* 3		TABLE ACCESS FULL	B	3	60
4		TABLE ACCESS FULL	A	4	80
-----					

Predicate Information (identified by operation id):

- 
- 1 - filter("F\_LEAK"("X")=1) <== This is correct
  - 2 - access("A"."ID"="B"."ID")
  - 3 - filter("B"."Y"<>'bbb')





# Toward v9.3, Beyond v9.3

- Feature list in v9.2
  - VIEWのsecurity\_barrier属性
  - FUNCTIONのleakproof属性
- Feature list in v9.3(?)
  - ALTER TABLE ... SET ROW SECURITY (...)
- Feature list in v9.4(???)
  - Writer side checks
  - Security label column
  - Label based mandatory row-level access control



# Syntax of Row-level Security (1/2)

```
ALTER <table_name>
    SET ROW SECURITY FOR <cmd>
    TO (<expression>);
<cmd> := ALL | SELECT | INSERT | UPDATE | DELETE
```

- <cmd> で指定したDML構文の実行時に、<expression>で指定した条件句(セキュリティポリシー)が付加される。
- クエリの条件句よりも、セキュリティポリシーが先に評価される。

```
postgres=> ALTER TABLE my_table
            SET ROW SECURITY FOR ALL TO (a % 2 = 0);
ALTER TABLE
postgres=> ALTER TABLE my_table
            SET ROW SECURITY FOR ALL TO
            (a = ANY(SELECT x FROM sub_tbl));
ALTER TABLE
```



## Syntax of Row-level Security (2/2)

```
ALTER t SET ROW SECURITY FOR ALL  
      TO (owner = current_user);
```

```
SELECT * FROM t WHERE f_leak(x);
```



security\_barrier  
付きサブクエリ

```
SELECT * FROM (  
      SELECT * FROM t WHERE owner = current_user  
) t WHERE f_leak(x)
```

- テーブルへの参照 → 条件句付きテーブルスキャンを含むサブクエリ (with security\_barrier) に置き換える。
- superuser の場合には適用されない。
  - ビルトインRLSの場合。extensionが勝手に付けるのは自由。



# How does RLS work? (1/2)

```
postgres=> ALTER TABLE t
           SET ROW SECURITY FOR ALL TO (owner = current_user);
ALTER TABLE
```

```
postgres=> EXPLAIN (costs off)
           SELECT * FROM t WHERE f_leak(b) AND a > 0;
           QUERY PLAN
```

```
-----
Subquery Scan on t
  Filter: f_leak(t.b)
    -> Index Scan using my_table_pkey on t t_1
        Index Cond: (owner = "current_user"())
        Filter: (a > 0)
(5 rows)
```



# How does RLS work? (2/2)

```
postgres=> EXPLAIN (costs off)
            UPDATE t SET b = b WHERE f_leak(b);
            QUERY PLAN
```

---

Update on t

```
-> Subquery Scan on t_1
    Filter: f_leak(t_1.b)
-> Index Scan using my_table_pkey on t t_2
    Index Cond: (owner = "current_user"())
```

```
postgres=> EXPLAIN(costs off)
            DELETE FROM t WHERE f_leak(b);
            QUERY PLAN
```

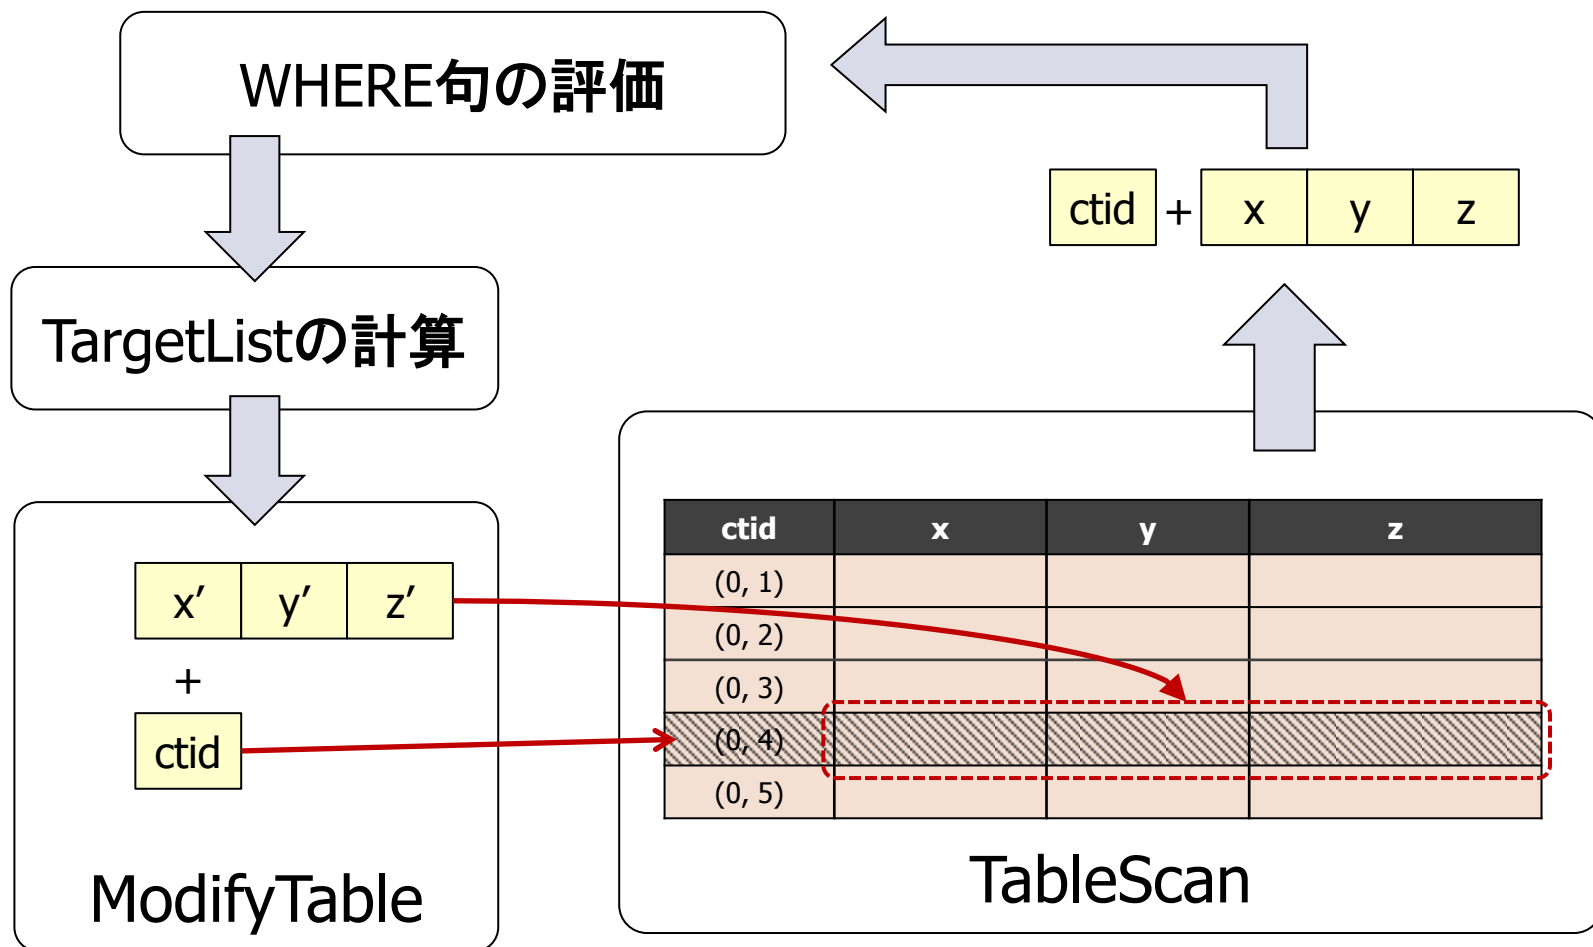
---

Delete on t

```
-> Subquery Scan on t_1
    Filter: f_leak(t_1.b)
-> Index Scan using my_table_pkey on t t_2
    Index Cond: (owner = "current_user"())
```

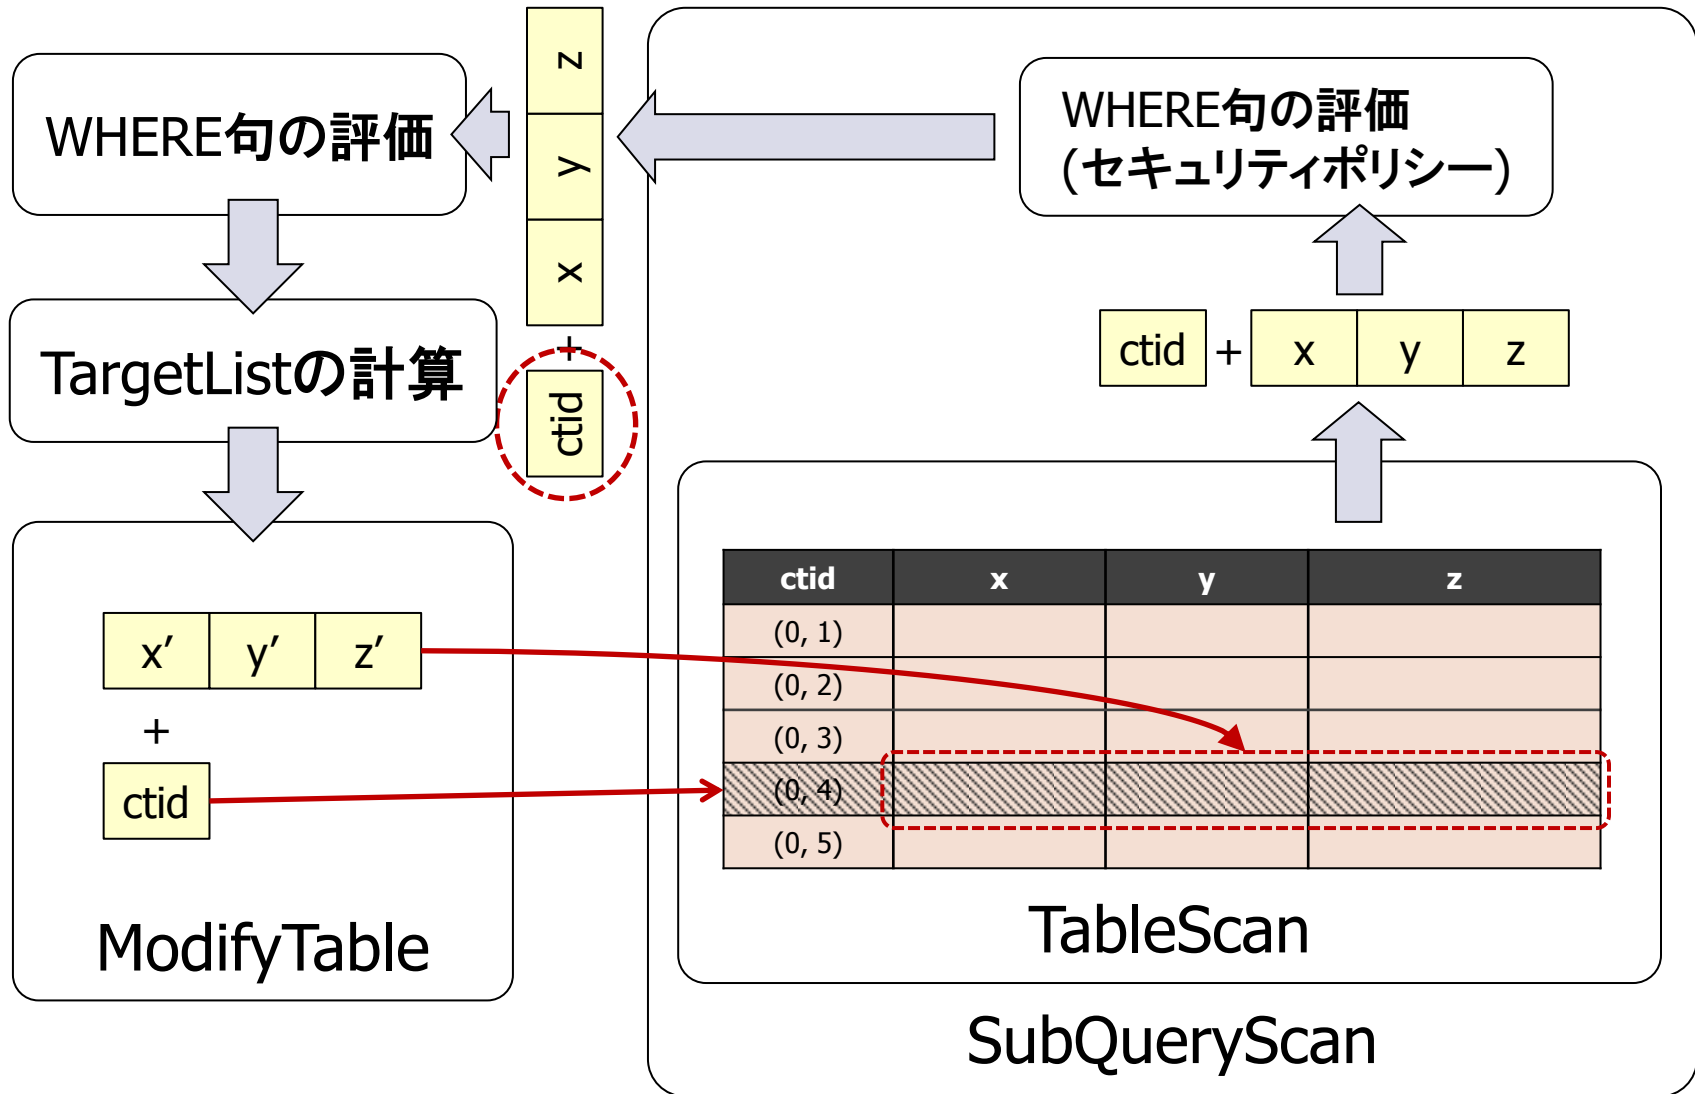


# Table Update and RLS (1/2)





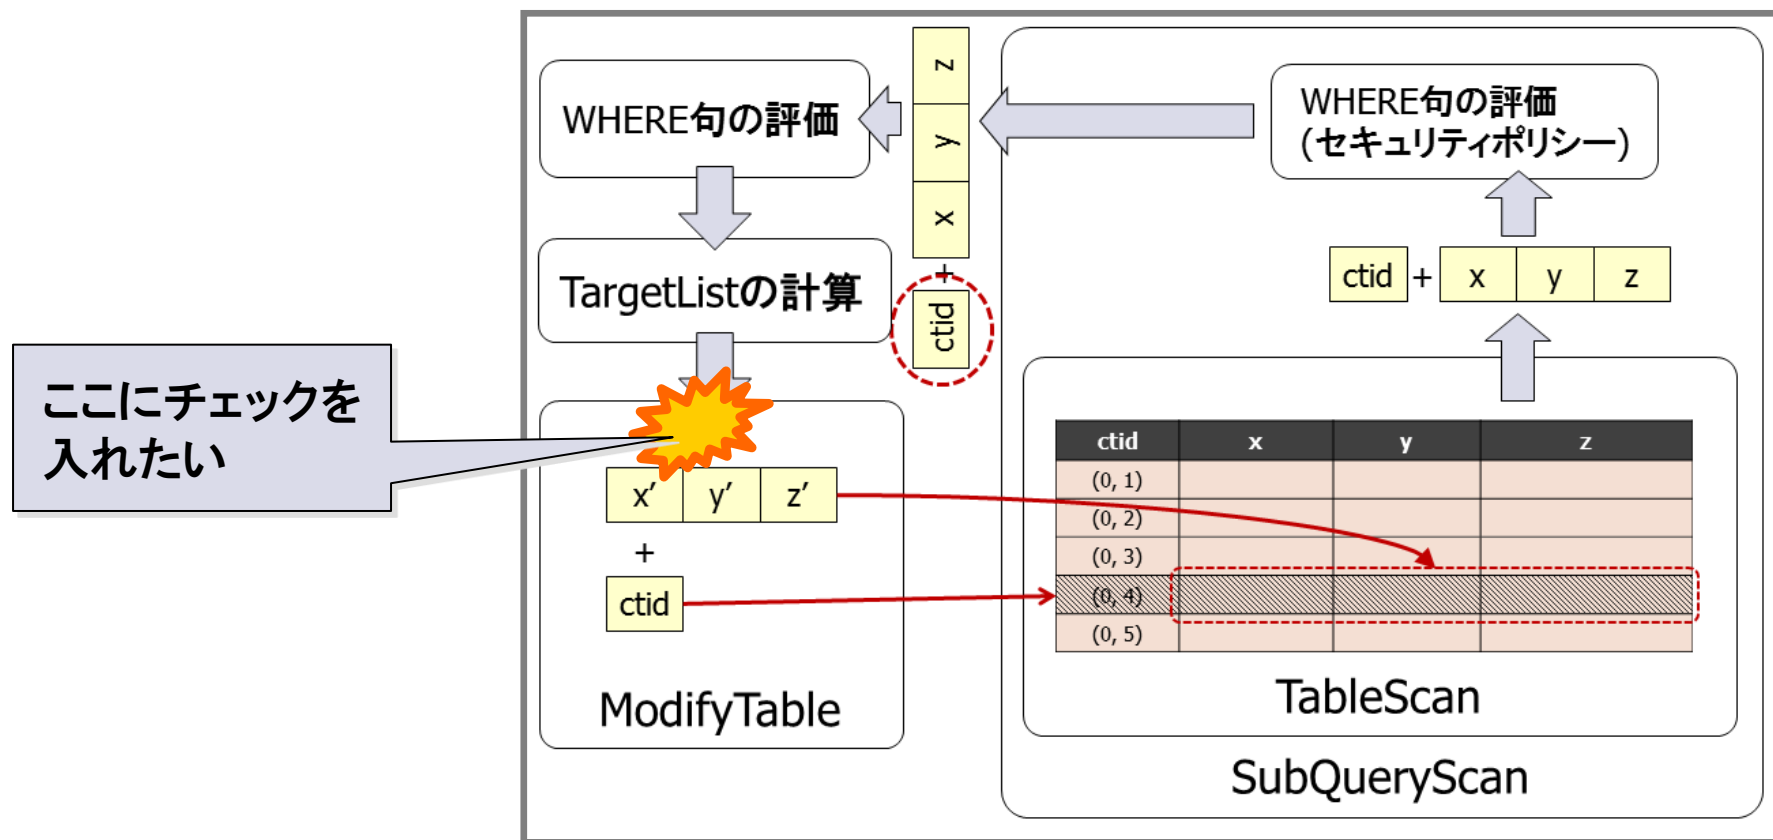
# Table Update and RLS (2/2)





# Future development (1/2)

- SQLコマンド毎に異なるセキュリティポリシーの設定
- INSERT/UPDATEの直前にもチェックを入れる
  - TargetListを計算した結果、それはPolicy Violationな値かもしれない!!







## Future development (2/2)

- Labal-based Row-level Security
  - 要は、SE-PostgreSQLの行レベルアクセス制御対応
- 実現に向けたプロセス
  - Row-level Security基本機能
  - INSERT/UPDATE直前のチェック機能
  - Security Label用の列を自動的に追加する機能
  - 要素の動的追加が可能な Enum 型
  - ↑ これらを全部活用して contrib/sepgsql の機能拡張😊



# Any Questions?